

Deep Q-Learning for High-Dimensional Control Tasks

Yoav Sabag

Yftach Gil



Agenda

- ▶ Introduction
- ▶ Methods
- ▶ Results
- ▶ Summary
- ▶ Conclusion

Introduction


The need for controlling devices

- ▶ Cameras are highly common high-dimension sensor
- ▶ Control application: Robots, Autonomous vehicles, Traffic optimization




The problem: how to perform control in a high-dimensional environment?


- ▶ Input: image (video) observation
- ▶ Output: action from a set of possible action
- ▶ Target: Human level or higher control for a task



Solution - Reinforcement learning




Methods




Environment - Cart Pole

- ▶ Action: 0=Push Left, 1=Push Right
- ▶ Goal: keep the pole upright for as long as possible
- ▶ Reward: +1 for each step
- ▶ The Observation Space is [Position, Velocity, Angle, Angular Velocity]
- ▶ Episode end conditions:
 - ▶ Pole Angle is greater than $\pm 12^\circ$
 - ▶ Cart Position is in the range ± 2.4
 - ▶ Time is “long enough”: steps reach 500 (for v1) or 200 (for v0)



Environment

- ▶ Cart Pole v1 from GYM
- ▶ <https://www.youtube.com/watch?v=B4E9tGmONn0>
- ▶ Observation space is converted from vector of size 4 to image (frame) of 90x40




Metrics and evaluation

Average reward: The average score obtained by the agent over 100 of episodes

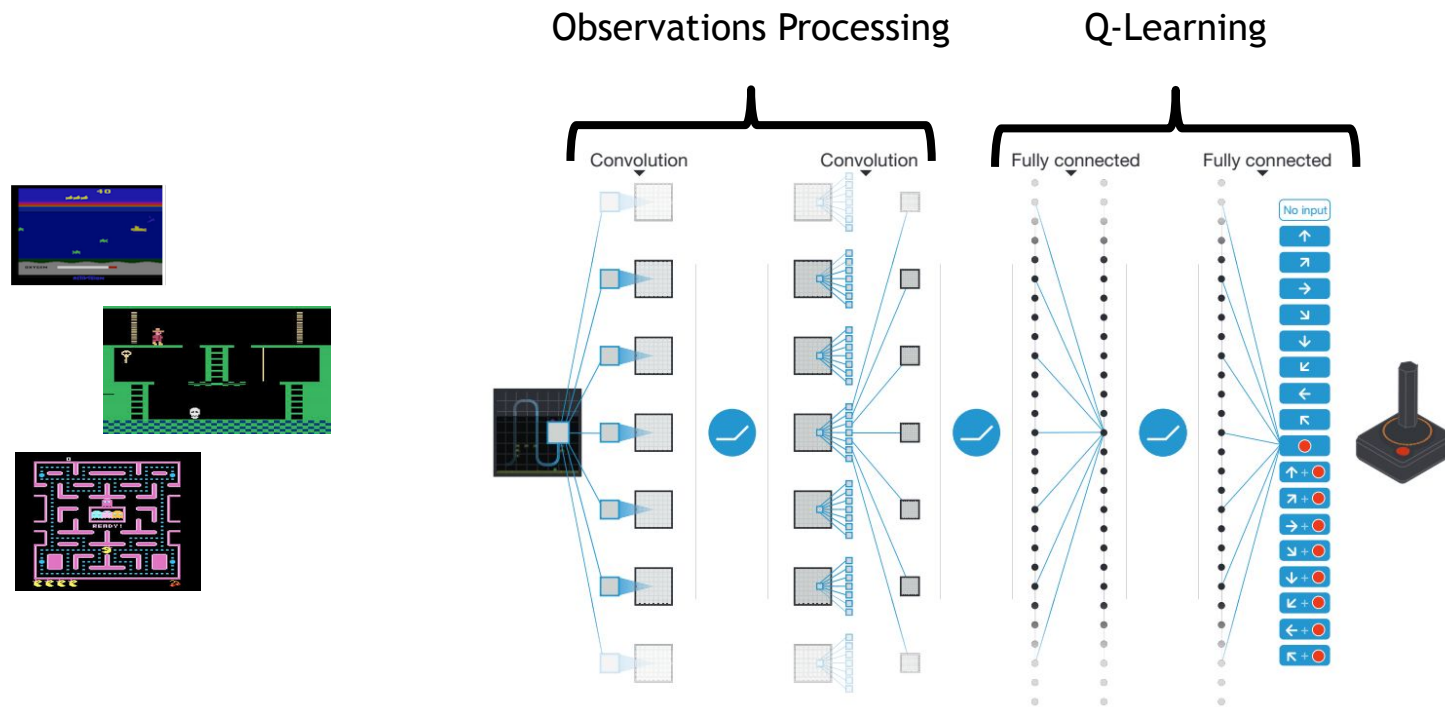
Episode length: the duration of each game episode (number of steps until end)

Training Episodes: the number of training episodes required for the agent to achieve defined level of performance.



Referenced Literature

- ▶ “Human-level control through deep reinforcement learning”
- ▶ RL model to master dozens of games with same model and hyperparameters
- ▶ Excel human level control over tasks
- ▶ Deal with high-dimensional sensory input



Why Deep Q-Learning

- ▶ **Efficient** Learns successful policies directly from high dimensional inputs
- ▶ **Handling Non-linear Policies:** masters complex nonlinear strategies
- ▶ **Robustness to Different Domains:** as demonstrated in its application to multiple Atari games
- ▶ **Integration of Reward and Perception:** like the human brain
- ▶ **Continuous Improvement:** keep learning through interactions



METHODS - Model parts

- ▶ Agent

- ▶ in charge of getting the frames, performs the actions determined by the policy

- ▶ DQN:

```
(conv1): Conv2d(3, 32, kernel_size=(8, 8), stride=(4, 4), padding=(1, 1))
(conv2): Conv2d(32, 64, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1))
(conv3): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(fc1): Linear(in_features=2816, out_features=32, bias=True)
(fc2): Linear(in_features=32, out_features=64, bias=True)
(fc3): Linear(in_features=64, out_features=128, bias=True)
(out): Linear(in_features=128, out_features=2, bias=True)
```

- ▶ Experience

- ▶ stores each step's state, action, next_state and reward

- ▶ Replay memory

- ▶ During training replay memories are queried from a batch to "replay" the agent's experience. Using replay memory the agent "remembers" best action in a defined number of episodes. This increases robustness (prevents degradation).

- ▶ Policies:

- ▶ policy_net and target_net policies are identical nets
 - ▶ the target_net is fed with "next_states" and then the loss is computed between policy_net current state and next state.

METHODS - Model Parts

- ▶ initialize env, create policy_net and target_net
- ▶ take action → reward, state
- ▶ store experience(=state, action, next_states, reward) in replay buffer
- ▶ sample random batch from memory
- ▶ get q-values for current-policy next-target
- ▶ calculate Q-values using Bellman Optimality Equation):

$$q_values = reward + \gamma \cdot next_q_values$$

- ▶ optimal action-value function:

$$Q^*(s,a) = \max_{\pi} \mathbb{E}[r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \dots | s_t = s, a_t = a, \pi]$$

- ▶ Loss:

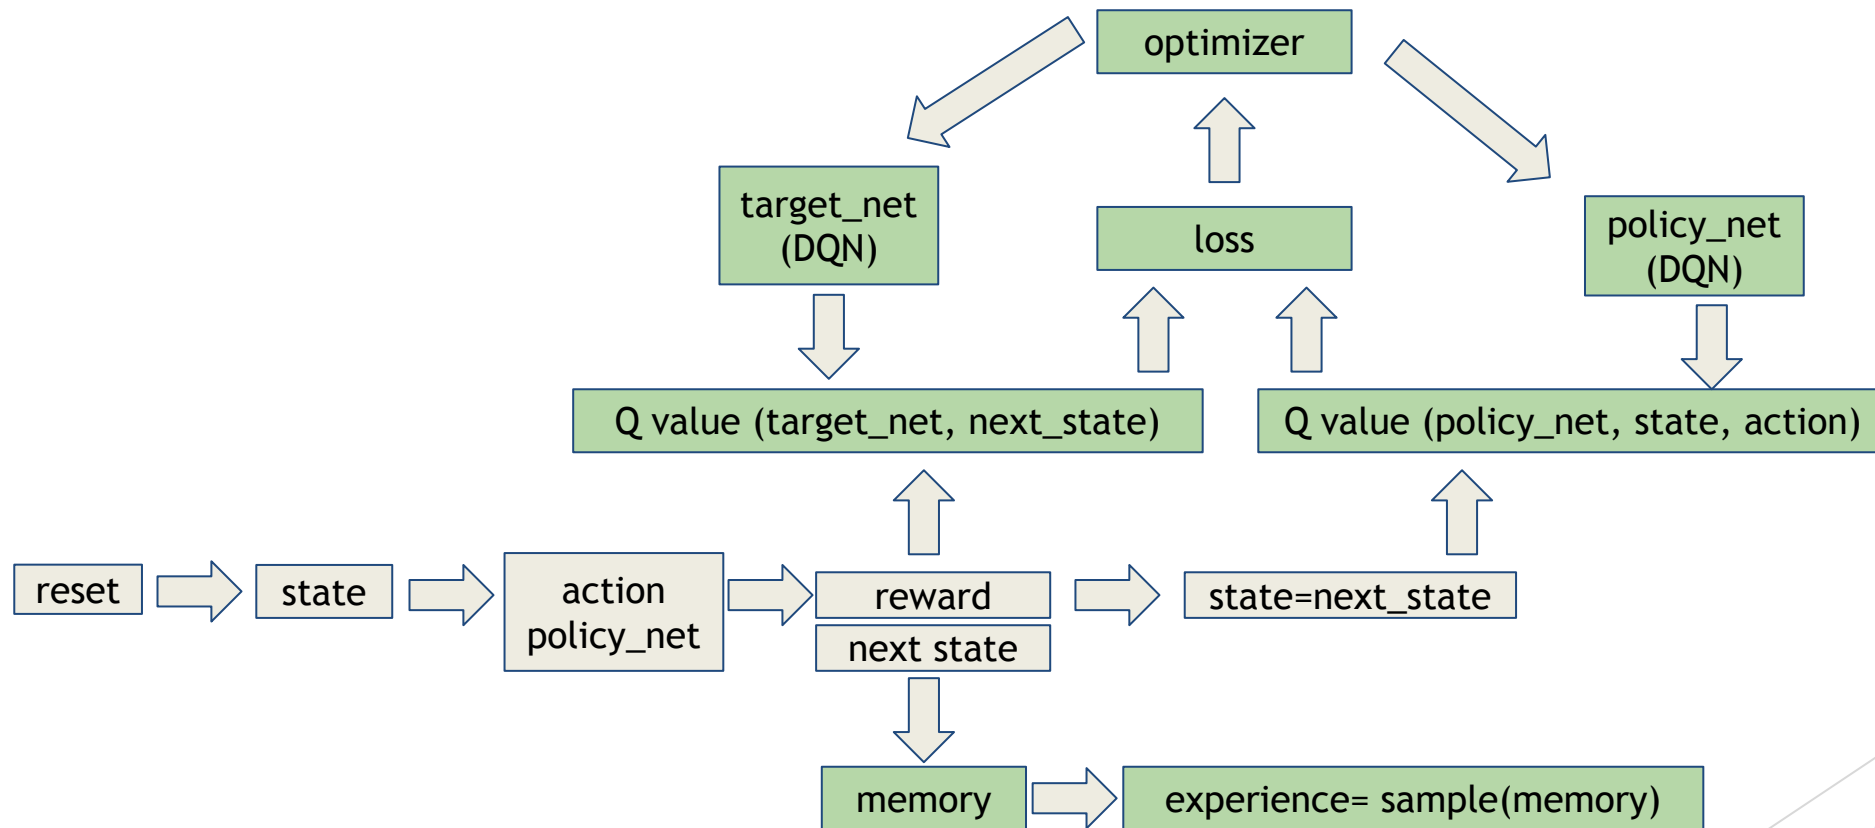
$$MSE = \frac{1}{n} \sum (current_q_value - target_q_value)^2$$

- ▶ optimal loss function:


$$L_i(\theta_i) = \mathbb{E}_{(s,a,r,s') \sim \mathcal{U}(D)} \left[\left(r + \gamma \max_{a'} Q(s', a'; \theta_i^-) - Q(s, a; \theta_i) \right)^2 \right]$$

- ▶ optimizer: update the network weights

METHODS - Architecture Scheme




Results



Link to play: <https://jeffjar.me/cartpole.html>


Training episode example of 167 steps



Experiment 1 Results (Reference 1)

DQN for cart-pole by TensorFlow

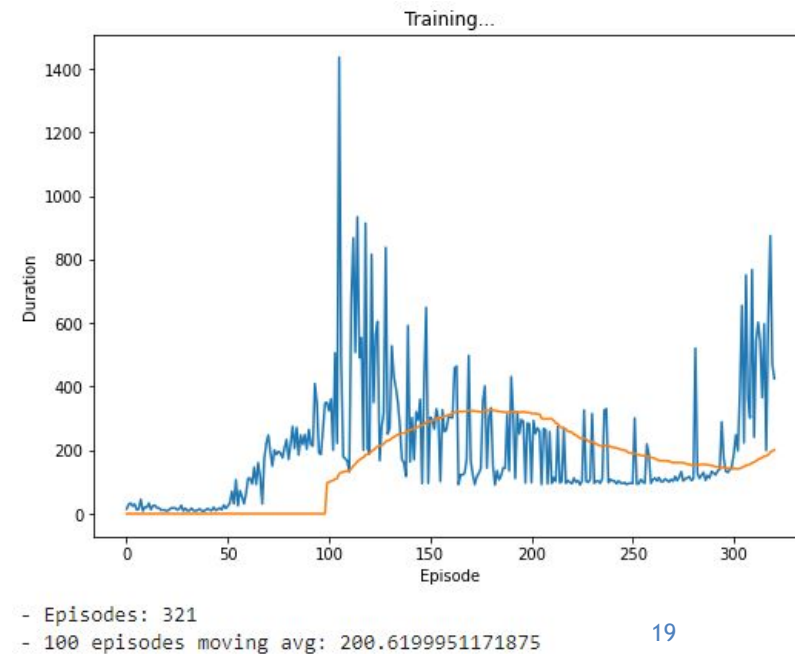
- ▶ observations as 4 dimensional vector [Position, Velocity, Angle, Angular Velocity]
- ▶ trained for 20,000 episodes (iterations)
- ▶ maximum reward is reached after 2500 episodes and stable after 10,000
- ▶ iteration = episode (200 steps maximum)
- ▶ number of steps = duration
- ▶ reward (=return), +1 for each step



Experiment 2 Results (Reference 2)

DQN - OpenAI Gym CartPole with PyTorch Kaggle notebook by dsxavier


- ▶ observations as vector (m=4, not a frame)
- ▶ trained for few hundreds episode 321
- ▶ average reward of 195 is reached 120 episodes




Main Paper (Reference 3)

“Human-level control through deep reinforcement learning”


- ▶ Control through deep reinforcement learning
- ▶ Metric: Q-Value, game score
 - ▶ Normalized performance of DQN score:
$$\text{Score} = 100 * (\text{DQN score} - \text{random play score}) / (\text{human score} - \text{random play score})$$
 - ▶ Training score: Game score



Each point is the average score achieved per episode after the agent is run with e-greedy policy ($\epsilon = 0.05$) for 520 k frames on Space Invaders game




Average score achieved per episode for Seaquest game



Experiment 3

- ▶ The average results achieved so far are 22.5% of the maximum possible

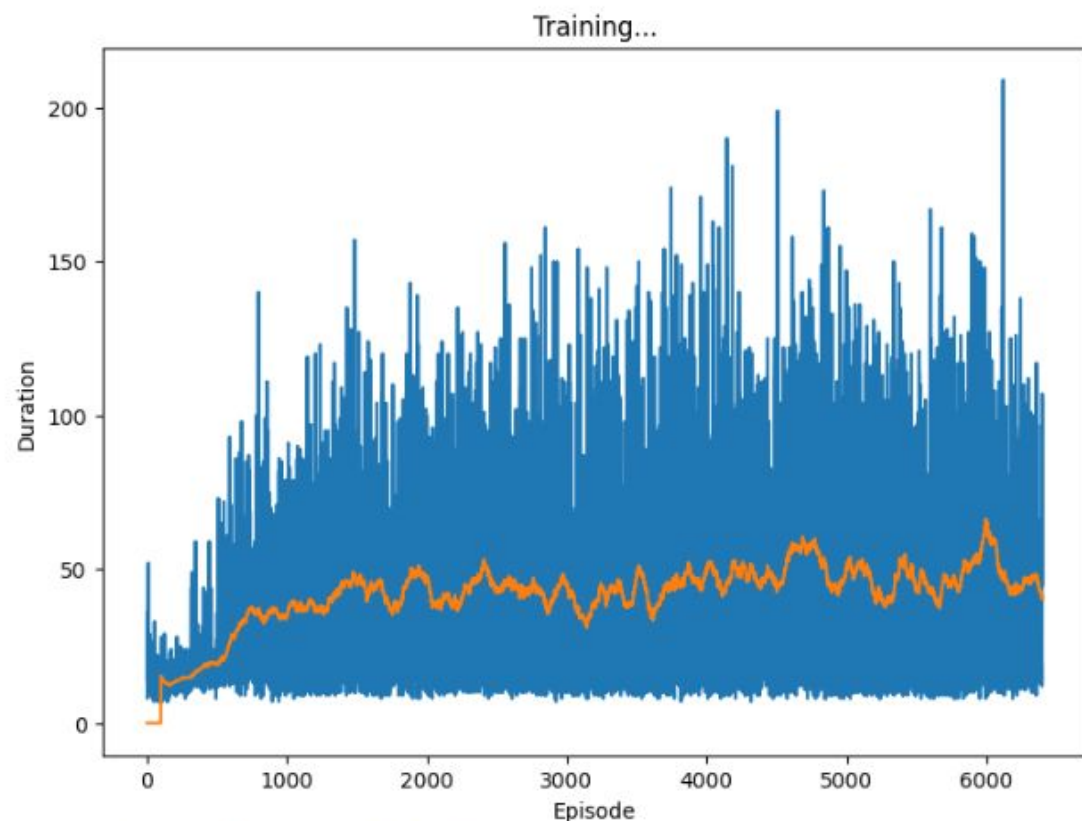


- Episode: 1394, Duration [steps]: 19
- 100 episodes moving avg: 44.97999954223633
- epsilon_rate: 0.01

```
DQN(  
  (conv1): Conv2d(3, 32, kernel_size=(8, 8), stride=(4, 4), padding=(1, 1))  
  (conv2): Conv2d(32, 64, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1))  
  (conv3): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
  (fc1): Linear(in_features=2816, out_features=32, bias=True)  
  (fc2): Linear(in_features=32, out_features=64, bias=True)  
  (fc3): Linear(in_features=64, out_features=128, bias=True)  
  (out): Linear(in_features=128, out_features=2, bias=True)  
)  
batch_size:      256  
gamma:           0.999  
eps_start:       1  
eps_end:         0.01  
eps_decay:       0.001  
target_update:   10  
memory_size:     100000  
lr:              0.001  
num_episodes:    768
```

Duration of training: 11:28

Experiment 3



- Episode: 6400, Duration [steps]: 45
- 100 episodes moving avg: 40.5
- epsilon_rate: 0.01

```
DQN(  
  (conv1): Conv2d(3, 32, kernel_size=(8, 8), stride=(4, 4), padding=(1, 1))  
  (conv2): Conv2d(32, 64, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1))  
  (conv3): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
  (fc1): Linear(in_features=2816, out_features=32, bias=True)  
  (fc2): Linear(in_features=32, out_features=64, bias=True)  
  (fc3): Linear(in_features=64, out_features=128, bias=True)  
  (out): Linear(in_features=128, out_features=2, bias=True)  
)  
batch_size:      64  
gamma:           0.999  
eps_start:       1  
eps_end:         0.01  
eps_decay:       0.001  
target_update:   10  
memory_size:     100000  
lr:              0.001  
num_episodes:    6400
```

Duration of training: 113:26

Summary

- ▶ **Customisation of the DQN Methodology:** Originally developed for Atari game environments, the Deep Q-Network (DQN) was converted for the cart-pole control task.
- ▶ **High-Dimensional Input:** The environment was modified to output video frames of 40x90 resolution
- ▶ **Agent:** Successfully received video frames and Convolution layers were added to the DQN in order to process the high dimension input.
- ▶ **Comparative Performance Analysis:** Despite these enhancements, the convolutional DQN achieved a maximum average reward of 44.9, which is 22.5% of the potential maximum for the Cart-Pole-v1 environment.
- ▶ **Benchmarking Against Reference Models:** Reference models demonstrated superior performance, achieving higher rewards and requiring fewer training episodes to converge.

Conclusion

- ▶ **Environment adaptation:** In this work a system designed to play atari games was modified for solving control task.
- ▶ **Challenges with High-Dimensional Complexity:** The solution for the control task achieved results obtained was lower than simpler methods based on lower dimensional environments.

Further improvements:

- ▶ The high dimensional complexity was challenging for the system and it didn't excel in comparison to the references of DQN without convolution.
- ▶ Hyper parameter fine tuning and Architecture improvement are needed
- ▶ More experiments are needed to be done on other control environments to determined if the model is robust across a range of control tasks.

Referenced and Related Work

1. Reference(1): tensorflow implementation of DQN for “cart pole” control environment
https://www.tensorflow.org/agents/tutorials/1_dqn_tutorial
2. Reference(2): dsxvier library for DQN
<https://www.kaggle.com/code/dsxavier/dqn-openai-gym-cartpole-with-pytorch>
3. Reference(3): main paper for the work
<https://web.stanford.edu/class/psych209/Readings/MnihEtAlHassibis15NatureControlDeepRL.pdf>
4. Reference(4): environment documentation
https://www.gymnasium.dev/environments/classic_control/cart_pole/
5. Reference(5): DQN with high dimensionality input over space invader environment
<https://www.kaggle.com/code/yaaryan/space-invaders-game-using-deep-q-networks>

